



*If at any point in time you need help on solving this assignment, view our demo video to understand the different steps of the code.*

## 1: Import the dataset

```
In [5]: #Import the diabetes dataset
df_diabetes = pd.read_csv('pima-indians-diabetes.data', header =None)
```

```
In [7]: #View the first five observations of the dataset
df_diabetes.head()
```

### 3: Find the features of the dataset

```
In [16]: #Use the .NAMES file to view and set the features of the dataset
df_diabetes.columns= ['pregnant','glucose','bp','skin','insulin','bmi','pedigree','age','label']
```

```
In [17]: #Use the feature names set earlier and fix it as the column headers of the dataset
df_diabetes.columns= ['pregnant','glucose','bp','skin' ,'insulin','bmi','pedigree' ,'age','label']
```

```
In [18]: #Verify if the dataset is updated with the new headers
df_diabetes.columns

Out[18]: Index(['pregnant', 'glucose', 'bp', 'skin', 'insulin', 'bmi', 'pedigree',
                'age', 'label'],
                dtype='object')
```

```
In [19]: #View the number of observations and features of the dataset
df_diabetes.shape
```

Out[19]: (768, 9)

#### 4: Find the response of the dataset

```
In [20]: #Select features from the dataset to create the model
feature_cols = ['pregnant', 'insulin', 'bmi', 'age']
```

```
In [21]: #Create the feature object
X = df_diabetes[feature_cols]
```

```
In [23]: #Create the reponse object
Y = df_diabetes[['label']]
```

```
In [24]: #View the shape of the feature object
X.shape
```

```
Out[24]: (768, 4)
```

```
In [25]: #View the shape of the target object
         Y.shape
```

```
Out[25]: (768, 1)
```

## 5: Use training and testing datasets to train the model

```
In [27]: #Split the dataset to test and train the model
from sklearn.model_selection import train_test_split
x_train , x_test , y_train ,y_test = train_test_split(X,Y,random_state=1)
```

## 6: Create a model to predict the diabetes outcome

```
In [29]: # Create a logistic regression model using the training set
from sklearn.linear_model import LogisticRegression
log = LogisticRegression().fit(x_train,y_train);

C:\Users\Mohannad\Anaconda3\lib\site-packages\sklearn\linear_model\logistic.py:432: FutureWarning: De
fault solver will be changed to 'lbfgs' in 0.22. Specify a solver to silence this warning.
  FutureWarning)

C:\Users\Mohannad\Anaconda3\lib\site-packages\sklearn\utils\validation.py:724: DataConversionWarning:
A column-vector y was passed when a 1d array was expected. Please change the shape of y to (n_sample
s, ), for example using ravel().
  y = column_or_1d(y, warn=True)
```

```
In [30]: #Make predictions using the testing set
log.predict(x_test)
```

[illegible]

```
In [34]: y_hat = log.predict(x_test)
```

## 7: Check the accuracy of the model

```
In [32]: #Evaluate the accuracy of your model
log.score(x_train,y_train)
```

```
Out[32]: 0.6857638888888888
```

```
In [33]: log.score(x_test, y_test)
```

```
Out[33]: 0.6979166666666666
```

```
In [37]: from sklearn import metrics
print( metrics.accuracy_score(y_test,y_hat) )

0.6979166666666666
```

```
In [63]: #Print the first 30 actual and predicted responses
print('the actual {}'.format(y_test.values[0:30]))
print(' the predicted {}'.format(y_hat[0:30]))

the actual [0 0 0 0 0 0 0 0 0 0 0 0 1 1 0 0 0 0 0 0 0 0 1 0 0 0 0 0 0 0]
the predicted [0 0 0 0 0 0 0 0 0 0 0 0 0 0 1 1 0 0 0 0 0 0 0 0 1 0 0 0 0 0 0]
```

In [ ]: