



Assignment 01: Evaluate the Ad Budget Dataset of XYZ Firm

The comments/sections provided are your cues to perform the assignment. You don't need to limit yourself to the number of rows/cells provided. You can add additional rows in each section to add more lines of code.

If at any point in time you need help on solving this assignment, view our demo video to understand the different steps of the code.

Happy coding!

1: Import the dataset

```
In [143]: #Import the required libraries
import numpy as np
import pandas as pd
```

```
In [144]: #Import the advertising dataset
advertising = pd.read_csv('Advertising Budget and Sales.csv',index_col= 0)
```

2: Analyze the dataset

```
In [145]: #View the initial few records of the dataset
advertising.head()
```

```
Out[145]:
```

	TV Ad Budget (\$)	Radio Ad Budget (\$)	Newspaper Ad Budget (\$)	Sales (\$)
1	230.1	37.8	69.2	22.1
2	44.5	39.3	45.1	10.4
3	17.2	45.9	69.3	9.3
4	151.5	41.3	58.5	18.5
5	180.8	10.8	58.4	12.9

```
In [146]: #Check the total number of elements in the dataset
print(advertising.size)
print(advertising.shape[0]*advertising.shape[1])
```

```
800
800
```

3: Find the features or media channels used by the firm

```
In [147]: #Check the number of observations (rows) and attributes (columns) in the dataset
advertising.shape
```

```
Out[147]: (200, 4)
```

```
In [148]: #View the names of each of the attributes
advertising.columns
```

```
Out[148]: Index(['TV Ad Budget ($)', 'Radio Ad Budget ($)', 'Newspaper Ad Budget ($)',
               'Sales ($)'],
              dtype='object')
```

4: Create objects to train and test the model; find the sales figures for each channel

```
In [149]: #Create a feature object from the columns
X = advertising.drop(columns='Sales ($)')
```

```
In [150]: #View the feature object
X.head()
```

```
Out[150]:
```

	TV Ad Budget (\$)	Radio Ad Budget (\$)	Newspaper Ad Budget (\$)
1	230.1	37.8	69.2
2	44.5	39.3	45.1
3	17.2	45.9	69.3
4	151.5	41.3	58.5
5	180.8	10.8	58.4

```
In [151]: #Create a target object (Hint: use the sales column as it is the response of the dataset)
Y = advertising[['Sales ($)']]
```

```
In [152]: #View the target object
Y.head()
```

```
Out[152]:
```

	Sales (\$)
1	22.1
2	10.4
3	9.3
4	18.5
5	12.9

```
In [153]: #Verify if all the observations have been captured in the feature object
X.shape
```

```
Out[153]: (200, 3)
```

```
In [154]: #Verify if all the observations have been captured in the target object
Y.shape
```

```
Out[154]: (200, 1)
```

5: Split the original dataset into training and testing datasets for the model

```
In [155]: #Split the dataset (by default, 75% is the training data and 25% is the testing data)
from sklearn.model_selection import train_test_split
x_train, x_test , y_train , y_test = train_test_split(X,Y,test_size=0.25, random_state=1)
```

```
In [156]: #Verify if the training and testing datasets are split correctly (Hint: use the shape() method)
print(x_train.shape , x_test.shape , y_train.shape , y_test.shape)
```

```
(150, 3) (50, 3) (150, 1) (50, 1)
```

6: Create a model to predict the sales outcome

```
In [157]: #Create a linear regression model
from sklearn.linear_model import LinearRegression
lr = LinearRegression().fit(x_train,y_train)
```

```
In [158]: #Print the intercept and coefficients
lr.intercept_ , lr.coef_
```

```
Out[158]: (array([2.87696662]), array([[0.04656457, 0.17915812, 0.00345046]]))
```

```
In [159]: #Predict the outcome for the testing dataset
lr.predict(x_test)
```

```
Out[159]: array([[21.70910292],
                [16.41055243],
                [ 7.60955058],
                [17.80769552],
                [18.6146359 ],
                [23.83573998],
                [16.32488681],
                [13.43225536],
                [ 9.17173403],
                [17.333853 ],
                [14.44479482],
                [ 9.83511973],
                [17.18797614],
                [16.73086831],
                [15.05529391],
                [15.61434433],
                [12.42541574],
                [17.17716376],
                [11.08827566],
                [18.00537501],
                [ 9.28438889],
                [12.98458458],
                [ 8.79950614],
                [10.42382499],
                [11.3846456 ],
                [14.98082512],
                [ 9.78853268],
                [19.39643187],
                [18.18099936],
                [17.12807566],
                [21.54670213],
                [14.69809481],
                [16.24641438],
                [12.32114579],
                [19.92422501],
                [15.32498602],
                [13.88726522],
                [10.03162255],
                [20.93105915],
                [ 7.44936831],
                [ 3.64695761],
                [ 7.22020178],
                [ 5.9962782 ],
                [18.43381853],
                [ 8.39408045],
                [14.08371047],
                [15.02195699],
                [20.35836418],
                [20.57036347],
                [19.60636679]])
```

```
In [160]: y_hat = lr.predict(x_test)
y = y_test
```

7: Calculate the Mean Square Error (MSE)

```
In [175]: #Import required libraries for calculating MSE (mean square error)
from sklearn import metrics
```

```
In [177]: #Calculate the MSE
mean_squared_error = np.sqrt(metrics.mean_squared_error(y,y_hat))
```

```
In [178]: print('the mean_squared_error is {}'.format(mean_squared_error))
```

```
the mean_squared_error is 1.404651423032895
```

```
In [ ]:
```